From input to final output

A user provides a Starling roleplay link or a finalized script. The workflow turns that source into reviewed, voiced node videos while preserving the story branches.

The complete flow



Repair loop: a failed image returns to image correction; a failed video returns to its source/prompt or EN2 render step, then is reviewed again. Only accepted videos enter stitching.

The following pages explain each stage, its skills and outputs. Technical references and storage details follow the plain-language flow. The 11 separate Word files contain the complete current SKILL.md text.

1 Confirm the inputs

The first step is to establish what the user wants to preserve and which style and voices to use.

Three mandatory questions

• **Source:** Is the input a Starling roleplay link or a finalized script document?

• **Style:** Use the actual visual references from the source roleplay, or provide a style image / Starling reference link for the finalized document.

• **Voices:** Which ElevenLabs voice ID or voice name should each speaking character and narrator use? Alternatively, explicitly select supplied final recordings or a silent production.

Two clear starting paths

Input	What happens
Starling link	Retrieve every available beat, choice, branch, recovery, consequence, ending, character and actual visual reference through Starling MCP.
Finalized document	Follow its words, speakers, scenes and routes directly. A style-reference roleplay supplies visual style and compatible beat context; it does not replace the finalized story.

Read-only source inspection may identify the complete speaker list before voice mapping is finalized. Building and media generation wait for the answers. Existing answers for the same source/version are reused; an ambiguous voice name is clarified rather than guessed.

Saved result: intake.json records the source mode, selected style, speaker list, voice choices and the actual user answers. Referenced local files have fingerprints so a later change can be detected.

Relevant skills: starling-production-master; starling-script-rebuilder; starling-audio-pack. Shared rule: pipeline/INTAKE.md.

2 Prepare and store the script

The script becomes the fixed reference for all later speech, visuals and branching.

Build from the complete source

For a Starling link, retrieve the full source through MCP and build a coverage list. Every source element is mapped to a scene, line, choice or route. Keep all characters and source choices; do not force the story into a fixed number of stops or a fixed duration.

For a finalized document, extract the content in reading order, including tables, and map it into the same structure without rewriting the approved prose. Ask about unclear speakers or missing route targets.

Keep the branches explicit

A node is one playable story unit. A Segment is a source-supported scene block; it may end with a decision or continue. Each choice points to its own destination. A same-Segment retry uses one explicit Try again button. Preserve any recovery route or ending already specified by the source.

One source for Word and JSON

The canonical script file is written once. The reusable builder validates it and produces a Word review document and a developer JSON graph from that same data. This keeps the dialogue and routes aligned. Word layout is rendered and inspected before a real script deliverable is approved.

Store the approved script in Google Docs

Confirm the connected account is darshan@en2.ai before writing. Save the real document ID and URL, then read back its native content and compare it with the canonical script. Google Docs is an additional script-storage gate; local files remain the production reference.

Outputs: canonical-script.json, source snapshots and coverage records, a review DOCX, matching developer JSON, and a native Google Doc with verified readback. The script builder emits script.docx and script.json in the selected output directory.

Relevant skills: starling-script-rebuilder; starling-production-master. Supporting integrations: Google Drive / Google Docs and Documents skills.

3 Create and accept the audio

Audio is completed before the final visual timing is planned.

Generate only the required speech

Inventory every spoken line with its speaker, language and branch. Use the selected ElevenLabs voices with the eleven_v3 model, or reuse explicitly selected final recordings. Choice buttons and other player controls are spoken only when the script says they are spoken.

Listen and measure

• Compare the entire recording with the script: all words, speaker identity, pronunciation, pauses and joins.

• Measure real duration and speech cues. A word-count estimate is not a final timing measurement.

• Save original recordings and actual provider job IDs. Reconcile an uncertain paid request before submitting again.

Keep narration and dialogue distinct

Type	How it is used
Narration	Kept as an unchanged audio master. It can continue over visual cuts and is added during final stitching.
Dialogue	Accepted recordings are supplied to EN2 with the clean scene, clear speaker assignments and measured turn order. Lip-sync is reviewed.
Speech-free coverage	Uses an explicit empty speech inventory for the interval. No dummy narration is generated.

A derived conversation track may join separate character recordings in the correct order. Any excerpt keeps its original speed and records where it came from. Replacing audio requires new timing checks for the affected visuals.

Outputs: accepted audio files, audio/audio_manifest.json, speaker/voice mappings, duration and cue records, listening evidence and any derived dialogue excerpts.

Relevant skills: starling-audio-pack; starling-production-master. Connection: ElevenLabs MCP.

4 Plan the visual coverage

Each visual explains a particular meaning or action in the script at the right point in the audio.

Map meaning to time

Read the whole story and all branches. Break each scene into meaningful visual beats: an introduction, location, problem, action, detail, reaction or outcome. Link each beat to its source line and actual accepted-audio interval. One script segment may need several storyboards.

Decide what the viewer must understand

• Record who and what is visible, the setting, current object state and the purpose of the shot.

• Separate a proposed action from one that has already happened. Keep the correct outcome for that branch.

• Use a wide view for orientation, a medium view for interaction and a detail view when a specific object matters.

• Include recognizable character introductions, relevant setting/task coverage, useful opening orientation and closing coverage for actual endings and completed Bad Outcomes.

Choose two related shots

Every storyboard contains Shot 1 and Shot 2. The second shot adds useful information, such as a detail or supported reaction. It is not simply an ending pose for the first image. Essential action is shown within a shot or additional coverage rather than implied only by switching poses.

Place the hard cut at a measured audio cue. It need not fall halfway through the clip. If there is too little time for both shots, revise the grouping of beats rather than create a rushed cut. Plan supported character movement or a subtle camera move within each shot.

Outputs: a full asset inventory, source-to-audio-to-shot map, image prompts, reference assignments, shot purposes, planned timing and director/cinematography review records.

Relevant skills: starling-prompt-pack; starling-director; starling-cinematography.

5 Generate the image pack

Foundational references are approved before the storyboard and dialogue images that depend on them.

Generate in dependency order

Character references → environments → props → storyboards and dialogue scenes. Use native ChatGPT image generation and actual selected style images. Reference sheets define identity and materials; extra views on a character sheet are not extra characters in a scene.

The storyboard has exactly two panels



One column and two stacked panels.

Each panel uses the output aspect.

Black-and-white line drawings.

SHOT, time, CAMERA and ACTION below.

Caption timing matches the saved plan.

Each picture has a separate white caption strip below it. Keep arrows and motion symbols out of the artwork. Save picture-area boxes separately from labels. The video presents one edge-to-edge colour scene at a time, connected by a hard cut.

Dialogue follows a separate image route

Dialogue uses one clean full-frame scene containing the actual participants, with approved dialogue audio for the later EN2 request. A mixed scene can contain narration and dialogue blocks without changing its canonical story node.

Outputs: durable image files, versions and reference mappings; visuals/image_manifest.json; actual dimensions and panel boxes; and review records for every asset, board and scene.

Relevant skills: starling-image-pack-generator; starling-panel-corrector; starling-director; starling-cinematography.

6 Check images before animation

The check is general physical plausibility, not a list of rules for one particular prop.

Check	Plain-language question
Scale and proportion	Do people, objects and surroundings have believable size relationships for the approved style and camera distance?
Perspective and depth	Do the camera view, distances and placement agree?
Separation and clearance	Do solid bodies remain distinct, with enough space for the pose or action?
Support and contact	Is a seated person supported? Are feet grounded and grips physically plausible?
Orientation and construction	Are front/back faces, attachments, joints and object parts correctly arranged?
Occlusion	Do nearer surfaces correctly hide what is behind them?
Continuity	Do identity, ownership, scale and state remain consistent across the two shots?

These checks catch a person merging into a chair or table, an implausible human-to-machine scale, or front-side content appearing on the back of a held object. A legitimately large machine or a justified perspective change is not automatically a defect.

Also check source meaning, anatomy, cast count, clothing, lighting, scene state and clean presentation. Inspect both panels at useful resolution. If either panel fails, repair the actual image using its approved references, then recheck the complete replacement.

Approval is version-specific. Record the inspected file fingerprint, panel/region, observation and correction. Both panels and their relationship must pass before an animation prompt is accepted.

Relevant skills: starling-panel-corrector; starling-cinematography; starling-director; starling-verification. Shared reference: spatial-fidelity.md.

7 Turn boards into animation

Approved images become specific instructions and inputs for one EN2 video request per storyboard.

Write the prompt from the actual approved board

The animation prompt identifies each reference, maps the top panel to Shot 1 and the bottom panel to Shot 2, and describes only the supported movement within each shot. It specifies the clip duration, exact hard-cut cue, planned camera movement and scene relationships that must remain consistent.

Upload the actual files and verify them

For a board, upload relevant character, environment and prop references, followed by the complete two-panel board. Save the returned upload names and fingerprints. Retrieve the uploaded image bytes and inspect them before submitting a paid render; an upload success message alone is not verification.

Request	Images and sound sent to EN2
Storyboard video	Relevant asset references plus the complete board. Original narration is kept outside generation.
Dialogue video	Clean scene plus accepted dialogue recordings, with explicit speaker/turn assignments and any useful extra references.

What one board produces

One storyboard normally produces one video containing two full-screen shots with one hard cut. For example, an illustrative eight-second clip could show Shot 1 for three seconds and Shot 2 for five. Actual timings come from the accepted audio; this example is not a timing preset.

Use EN2 reference-to-video and Fast mode. Narration characters remain silent. Dialogue lip-sync belongs only to the active speaker. The prompt is a request to the model; the result still needs complete review.

Outputs: exact animation prompts and payloads, approved render plans, uploaded-file evidence, shot/cut timing and source-to-clip mappings.

Relevant skills: starling-animation-pack, with starling-director and starling-cinematography. Connection: EN2 video MCP.

Final narration presentation

Inspect the actual storyboard first, then write affirmative instructions that describe the intended result.

Two related shots with clear action

Use black-and-white 1x2 storyboards with SHOT, time, CAMERA and ACTION labels below each picture. Actions and labels must match the inspected staging. A quiet shot may use a subtle 3-5 percent push-in or a small pan while object states stay fixed. Preserve time of day, lighting, weather, identity and hairstyle.

Only relevant references in a fixed order

Attach character, environment, prop and storyboard references in that order. Omit categories not present in the storyboard. Asset sheets supply appearance; the storyboard controls visible cast, objects, scale and placement. Choose 2D, 2.5D or 3D animation from the approved colour references.

Affirmative prompt wording

> Render each illustrated scene individually as an edge-to-edge, full-frame colour shot. At the specified time, switch directly to the next fully formed shot using a hard cut. Treat the storyboard layout and caption strips exclusively as production guidance.

Review the finished video

Check the full result for the two timed scenes and direct hard cut. Fades, dissolves, overlays, mixed panels, white borders or frames, labels and storyboard layout appearing in the video are defects. Keep these defect lists in verification; generation prompts describe the desired full-frame result positively.

Relevant skills: starling-animation-pack; starling-verification. Final narration contract: references/narration-pairs.md. Worked prompt example: references/animation-template.md.

Animation prompt structure

The five-section template follows the supplied Claude prompt pack, adapted to the approved two-shot and audio workflow.

Prompt section	What it says
subject_definitions	Identify each attached Picture or Audio by its actual index. Explain its role and the shots where it applies.
summary	Explain the scene purpose, duration, exact visible cast and two-shot progression.
retention_analysis	For each reference and applicable shot, specify what is fully or partly preserved and explain any approved differences.
detailed_description	Describe each pictured starting pose, specific action, gaze/expression change and end hold. State the planned camera move and measured hard cut.
overall_soundscape	State the planned sound. Narration stays separate; dialogue follows accepted recordings. Music decisions stay in this section.

Make every reference unambiguous

Attach only relevant character sheets, environment and prop references, then the complete board. Save the ordered filenames, Picture indexes, roles and applicable shots. The board controls visible contents; a character sheet supplies identity, not additional people. Reconcile all indexes after upload.

Prevent additions before submitting

Record the people, objects, positions, contact, lighting and empty background visible in each panel. Compare every person, object and movement in the final prompt with that inventory. Fix incompatible staging in the image first. Save this comparison as shot_inventory_match evidence tied to the actual image files.

Example of precise action: Exactly one person is visible, seated at frame right with the phone already in his hand. He lowers his gaze to its dark screen, blinks once and holds. His forearm stays supported on his thigh; the empty passage retains its pictured state. Camera fixed.

The runtime requires five headings and the recorded inventory review; it cannot judge the pictures itself. Older animation approvals missing the new review or using the former format become stale before further submissions. Existing jobs and media remain preserved. Inspect the whole resulting video for added, duplicated or missing subjects, including during cuts.

Relevant skills: starling-animation-pack and starling-verification. Full reusable template: starling-animation-pack/references/animation-template.md.

8 Submit collect and repair

All approved work can be submitted before reviewing the generated output, while duplicate protections remain in place.

Submit all ready clips once

Submit the complete approved set in an ordered pass, respecting provider capacity. Save a local intent before each tool call and the real provider job ID immediately afterward. If a call has an uncertain outcome, reconcile the known job instead of making another paid request blindly.

Collect and account for every clip

Poll known IDs during the active session. Save returned states and download completed videos to durable storage. Record unfinished jobs and explicit failures separately. A missing item in a partial response is not automatically a failed render. A download failure triggers another download, not another video generation.

Review complete videos

Once the initial submission pass is accounted for, inspect every available complete picture and sound. Check the planned two shots, cut timing, planned framing, source meaning, physical relationships and motion throughout. Dialogue also requires complete speech, correct speakers and synchronization. Pending jobs remain visible blockers to final completion.

Repair through EN2 and inspect again

After an initial visual failure, redesign the scene once with simpler source-faithful staging, inspect fresh images and submit one EN2 Fast replacement. If it also fails, stop for review. Preserve the logical clip ID, originals and all paid-job IDs.

No still-image animation fallback. Failed animation is not replaced with a slow zoom or still sequence. FFmpeg is retained for checking, trimming, stitching, original-audio handling and export.

Relevant skills: starling-animation-monitor; starling-animation-pack; starling-verification; director/cinematography for their review checks.

9 Stitch and export

Only accepted EN2 videos are selected for the finished story nodes.

Join pictures in the script order

The assembly plan names each story node, its ordered clips, exact boundaries and original audio units. The code checks that all required clips are present, approved and unchanged. A clip-ID mismatch stops the handoff so footage is not silently skipped.

Add the original accepted audio

FFmpeg joins verified clips in canonical order. Narration uses accepted external speech; dialogue retains the selected video's embedded audio. Its generation route still uses clean scenes, selected voices, accepted speech and lip-sync checks.

Output is 1280 x 720 at 30 fps, scaled to fill and centrally cropped; inspect edge content. Assembly blocks an interval shorter than the selected dialogue. Extend that interval and shift following clips/audio units before retrying. Audio units must align with dialogue boundaries. Complete listening still establishes that the generated source contains every word. Audio is encoded once; extra music mixing and arbitrary source-in offsets remain unsupported.

Review the assembled result

• Check complete picture and sound, every spoken word and pause, dialogue synchronization and the joins between clips.

• Check every required branch. A correct-path review does not approve the alternative outcomes.

• Keep mutually exclusive outcomes in separate node videos. Player choices and retries remain in the script graph.

Final package: separate node MP4s, canonical script with choices, assembly plan, final review records, selected-output provenance and a README. Packaging does not publish an external player.

Relevant skills: starling-animation-pack; starling-verification. Local implementation: pipeline/postproduction.py and FFmpeg.

10 Where outputs are stored

The workspace keeps each production in its own project directory. The names below describe the current implementation and its recommended organization.

Workspace: C:\Users\darsh\Downloads\roleplay
Project root: projects/<project-name>/

```
projects/<project-name>/
  intake.json
  canonical-script.json
  pipeline.sqlite3
  source/                 [recommended source snapshots]
  script/                 [chosen builder output folder]
    script.docx
    script.json
  audio/
    audio_manifest.json
    ...original files, excerpts and cues
  visuals/
    image_manifest.json
    ...assets, boards, scenes and image reviews
  animation/
    animation-manifest.json
    assembly-plan.json
    ...prompts, payloads and upload evidence
  postproduction/
    state.sqlite3
    automation.json
    download-handoff.json
    raw/<provider-id>.mp4
    review/<clip-id>-<hash>/review.json
    edited/<node-id>-<version>/<node-id>.mp4
    edited/<node-id>-<version>/review.json
    delivery-manifest.json
    <project-name>-final.zip
```

Folders marked recommended or chosen are organization conventions, not automatically created source-retrieval folders. Manifest and postproduction filenames shown above are used by the current code. Existing productions may retain additional version folders.

Google Docs stores a native script copy remotely. Its actual returned ID/URL and a local readback record are preserved with the project. Media is downloaded locally rather than relying only on temporary provider URLs.

11 What each saved record means

Most records are small files that let the workflow resume accurately without repeating paid work.

Saved item	Purpose
intake.json	Source mode, style and voice answers, selected speaker mappings and supporting evidence.
canonical-script.json	The approved words and story graph used by all production stages.
audio/audio_manifest.json	Which recording belongs to which line/speaker, its duration and measured source cues.
visuals/image_manifest.json	Approved assets and images, versions, panel geometry, references and review information.
animation/animation-manifest.json	Clip inventory, approved inputs, content type, used duration and explicit runtime clip mapping.
animation/assembly-plan.json	Ordered clips and original audio units for each final node.
pipeline.sqlite3	Durable stage scopes, acceptance receipts, render intents, provider IDs and event history.
postproduction/state.sqlite3	Downloaded outputs, reviewed selections and assembled-node versions.
Review files	Who inspected which exact version, what was checked, observations and verdicts.
delivery-manifest.json	Which accepted outputs were selected for the final package.

A file fingerprint (SHA-256) detects changed bytes. If a source or accepted file changes, affected approval must be revisited. A fingerprint does not prove that the image or video is correct.

Provider responses and original media are preserved. Updating a skill does not rewrite an already submitted job or automatically migrate an older project into a new clip layout.

12 Connections and local tools

MCP is the connection that lets the active assistant call an external service. It is not a background process by itself.

Connection or tool	What it does	Where the result lives
Starling MCP	Retrieves the selected roleplay structure and available visual references.	Original source remains in Starling; snapshots and downloaded references are local.
ElevenLabs MCP	Resolves selected voices and generates speech with eleven_v3.	Real job IDs and downloaded audio stay with the project.
Native ChatGPT images	Generates/edits foundational assets, storyboards and dialogue scenes.	Actual images and review records are saved locally.
EN2 video MCP	Uploads references, submits Fast reference-video jobs and retrieves job states/results.	Real job IDs are recorded; video files are downloaded locally.
Google Drive / Google Docs	Stores a native script copy under the confirmed darshan@en2.ai connection.	Native document in Google Docs; real ID/URL and readback stored locally.
FFmpeg and ffprobe	Check media, normalize/trim clips, join audio/video and export final nodes.	Local files under postproduction.
Python and SQLite	Validate records, preserve progress and create the next handoff.	Local code and project databases.

The workflow uses these selected providers. It does not silently switch image/video generation to another service. Live tool names, schemas and limits must be checked when production is executed.

This report describes the implemented integration design; it is not a new live connectivity or account-access test. Native image generation and local FFmpeg are distinct from the external MCP services.

13 How the workflow is built

The system has three layers: written specialist instructions, saved project state and connected execution.

Skills explain the work

A SKILL.md tells the assistant what a specialist stage owns, which inputs it needs and what must be checked. Shared references hold rules used by several stages, such as intake, two-shot storyboard geometry and physical/spatial review. This keeps the individual skills short and consistent.

Runtime records enforce the handoffs

pipeline/runtime.py records the expected inventory for each stage and accepts a receipt only when its required evidence and file fingerprints are present. It saves an intent before a video request and the returned provider ID afterward. This helps avoid duplicate paid submissions and detects changed source files.

pipeline/postproduction.py connects scoped jobs to the assembly plan, records downloaded clips and reviews, selects accepted EN2 outputs, assembles nodes and checks final-package gates. It calls local FFmpeg; it does not run MCP calls.

The active session performs connected actions

The assistant reads the relevant skill, uses the connected tool, saves the actual result and records evidence. A Python command can issue the next handoff, but cannot independently fetch from the assistant’s MCP connections or claim that a video has been watched.

Term	Simple meaning
Scope	The exact list of items expected in a stage.
Receipt	A saved approval record tied to exact files and checks.
Manifest	An inventory and mapping of related production files.
Intent / provider ID	The local record before submission / the service’s actual returned job identifier.
Verified	Required review is complete for that exact version and scope.

14 Skill map and supporting files

Each separate Word document contains the complete current source of the named workflow skill.

Stage or responsibility	Skill document
Overall orchestration and mandatory intake	starling-production-master.docx
Source retrieval and canonical script	starling-script-rebuilder.docx
Selected voices, speech and measured cues	starling-audio-pack.docx
Beat mapping and image prompts	starling-prompt-pack.docx
Meaning, staging and action continuity	starling-director.docx
Framing, geometry and spatial relationships	starling-cinematography.docx
Image generation and durable outputs	starling-image-pack-generator.docx
Detailed still/panel checking and repair	starling-panel-corrector.docx
Animation prompts, EN2 repairs and stitching	starling-animation-pack.docx
Job reconciliation and output collection	starling-animation-monitor.docx
Evidence-based still/video/final review	starling-verification.docx

Additional files that support those skills

pipeline/INTAKE.md defines the required questions. WORKFLOW.md and POSTPRODUCTION.md explain operation and acceptance. narration-pairs.md and visual-routing.md define the two-shot route. spatial-fidelity.md defines general physical checks. scene-preservation.md and en2-submission.md govern animation handoff. handoff-contract.md and review-record.md define saved evidence.

The installed Script Rebuilder also includes references/data-contract.md, references/intake.md and scripts/build_script.py. The script skill is installed under the personal skill catalog; the other ten workflow skills live in the roleplay workspace.

15 Checks limits and final delivery

The workflow has implemented safeguards, but saved instructions alone do not make production unattended.

What has been tested

The current workflow passed 42 offline pipeline, intake and postproduction tests. Coverage includes narration reference routing, preserved historical records, duplicate prevention, review gates, embedded-dialogue audio, rejection of short dialogue intervals and real synthetic mixed-audio assembly. These results do not approve production media.

What those tests do not prove

Synthetic tests do not certify pronunciation, character identity, believable movement or lip-sync in real production. Complete listening and picture/sound inspection remain required. The standard skill validator lacked PyYAML in the update environment, so a separate structural audit checked frontmatter and links.

Current operational limits

• No persistent background runner is active. Collection through MCP needs an active session unless a supported runner is separately demonstrated.

• Full audiovisual review capability must be established for the actual files. A completed render or contact sheet is not enough.

• Older project manifests and clip mappings need explicit reconciliation before adopting the new workflow. Existing media and paid IDs remain preserved.

• The assembler uses external narration and embedded dialogue audio at 720p/30 fps. It does not mix additional ambience/music.

Delivery is complete only when: all required nodes and branches are present, their selected EN2 clips and final picture/sound pass review, original speech is intact, joins are checked, runtime receipts pass and the final ZIP is verified.

This PDF explains the current workflow as recorded on 17 October 2026. Its companion Word files reproduce the current skill instructions for manual checking. No new production run, publication or live monitoring is implied.